Prac Exercise 02
PostgreSQL: a client-server RDBMS

Aims

This exercise aims to get you to:

- install your PostgreSQL database server at CSE
- create, populate and examine a very small database

You ought to get it done by end of Week 2, since you'll need it to start working on Assignment 1.

Background

PostgreSQL runs on most platforms, but the installation instructions will be different for each one. There are downloadable binaries for many platforms, but they place restrictions on how your server is configured. Other installations come with a GUI front-end; I think this moves you too far away from the server, and simply gets in the way ... but maybe that's just me.

On Linux and MacOS, I have found it best to install from source code; the installation is relatively straightforward and you can specify precisely where on your filesystem you want the PostgreSQL files to be located. If you don't want to run PostgreSQL on your own machine, this Prac Exercise describes how to set it up on your CSE account. You can still work from home by accessing your server through VLab or `ssh`.

If you *do* want to install PostgreSQL on your home machine, you'll need to work out how to do that yourself. There are plenty of online resources describing how to do this for different operating systems (type "install postgresql" at Google). It probably doesn't matter if you use a slightly different version at home, since we'll be using a subset of SQL and PLpgSQL that hasn't changed for a while; any version around 11/12/13/14 is ok for the prac work in this course.

The practical work for the assignments can be carried out on a CSE machine called `d.cse.unsw.edu.au`. You run your own PostgreSQL server on this machine and are effectively the database administrator of this server. This machine has been configured to run large numbers* of PostgreSQL servers.

> * Note: "large numbers" is around 300. If you leave your work to the last minute, and find 400 other students all trying to run PostgreSQL on `d.cse`, performance will be sub-optimal. Of course, you can avoid this potential bottleneck by installing and running PostgreSQL on your home machine.

You should *not* use other CSE servers such as `wagner` and `weill` for running PostgreSQL; if you do, your PostgreSQL server will most likely be terminated automatically not long after it starts.

> Reminder: You should *always* test your work on the CSE machines before you submit assignments, since that's where we'll be running our tests to award your marks.

For brevity, we'll refer to "`d.cse.unsw.edu.au`" as "`d.cse`" in the rest of this document.

What `d.cse` provides is a large amount of storage and compute power that is useful for students studying databases. You should have access to `d.cse` because you're enrolled in COMP3311. You can access `d.cse` using either `ssh` or VLab.

In the examples below, we have used the `$` sign to represent the prompt from the command interpreter (shell). In fact, the prompt may look quite different on your machine (e.g., it may contain the name of the machine you're using, or your username, or the current directory name). All of the things that the computer types are in `this font`. The commands that **you** are supposed to type are in **`this bold font`**. Some commands use `$USER` as a placeholder for your CSE username.

Exercises

Stage 1: Getting started on d.cse.unsw.edu.au

Log in to `d.cse.unsw.edu.au`. If you're not logged into `d.cse` nothing that follows will work.

You can log into `d.cse` from a command-line (shell) window on any CSE machine (including `vlab`) via the command:

```
$ ssh d.cse.unsw.edu.au
```

If you're doing this exercise from home, you can use any `ssh` client, but you'll need to refer to `d.cse` via its fully-qualified name:

```
$ ssh YourZID@d.cse.unsw.edu.au
```

From home, an alternative is to use VLab. This requires a VNC client (e.g. TigerVNC). Use the VNC server

```
d.cse.unsw.edu.au:5920
```

You can check whether you're actually logged in to `d.cse` by using the command:

```
$ hostname
nw-syd-vxdb
```

Your home directory at CSE is directly accessible from `d.cse`. Run the `ls` command to check that you are indeed in your CSE home directory.

The first time you log in to `d.cse`, it automatically creates a directory to hold your databases:

```
$ ls -al /localstorage/$USER
```

This directory is initially empty, but we're about to put the files for a PostgreSQL server into it.

Stage 2: Setting up your environment

PostgreSQL needs certain configuration setting the be just right bfore it will work. Part of setting up this environment involves setting some shell environment variables.

The following commands will set up the environment appropriately:

```
I=/usr/lib/postgresql/13
PGDATA=/localstorage/$USER/pgsql/data
PGHOST=$PGDATA
LD_LIBRARY_PATH=$I/lib
PATH=$I/bin:$PATH
export PGDATA PGHOST LD_LIBRARY_PATH PATH
```

```
alias p0="$I/bin/pg_ctl stop"
alias p1="$I/bin/pg_ctl -l $PGDATA/log start"
```

The critical environment variables here are `PGDATA` (which indicates where all of the PostgreSQL data files are located), and `PGHOST` (which tells where the sockets are located that you use to connect to the PostgreSQL server). Of course, all of the other settings are important as well.

A useful place for these commands is in a file called

```
/localstorage/$USER/env
```

You should create this file and copy the above commands into it.

Once you've done this, running the command

```
$ source /localstorage/$USER/env
```

will set the environment appropriately.

You will need to set your environment each time you login to `d.cse` for a session with PostgreSQL.

If you know how to write shell scripts, you could modify your `.bashrc` so that it `source`'d the `env` script automatically, each time you login to `d.cse`.

Stage 3: Setting up your PostgreSQL Server

A necessary first step to installing a PostgreSQL server is to set up your environment as described in Stage 2. Once you've done that, you can create the directories and files to manage your databases.

The `initdb` command creates these directories and places some configuration files in them. You only need to run `initdb` once (unless you need to completely reinstall your PostgreSQL server from scratch).

If you used the environment setting described above, `initdb` will create your PostgreSQL server setup in the directory:

```
/localstorage/$USER/pgsql
```

An example of using `initdb`

```
$ which initdb
/usr/lib/postgresql/13/bin/initdb
$ initdb
The files belonging to this database system will be owned by user "jas".
This user must also own the server process.

The database cluster will be initialized with locale "C.UTF-8".
The default database encoding has accordingly been set to "UTF8".
The default text search configuration will be set to "english".

Data page checksums are disabled.

creating directory /localstorage/jas/pgsql/data ... ok
creating subdirectories ... ok
selecting dynamic shared memory implementation ... posix
selecting default max_connections ... 100
```

```
selecting default shared_buffers ... 128MB
selecting default time zone ... Australia/Sydney
creating configuration files ... ok
running bootstrap script ... ok
performing post-bootstrap initialization ... ok
syncing data to disk ... ok

initdb: warning: enabling "trust" authentication for local connections
You can change this by editing pg_hba.conf or using the option -A, or
--auth-local and --auth-host, the next time you run initdb.

Success. You can now start the database server using:

    pg_ctl -D /localstorage/jas/pgsql/data -l logfile start
```

Naturally, the occurrences of `jas` will be replaced by your zID.

Also, don't get too excited and run the above `pg_ctl` command just yet. There is a bit more setup to do ...

```
$ cd /localstorage/$USER/pgsql/data
$ edit postgresql.conf
```

where you replace the word *edit* by your favourite text editor.

The `postgresql.conf` file is the main configuration file for your PostgreSQL server. You need to modify it so that it works with the environment you set up above. This requires changes to the lines highlighted in red in this excerpt from `postgresql.conf`

```
...
# - Connection Settings -

#listen_addresses = 'localhost'      # what IP address(es) to listen on;
                    # comma-separated list of addresses;
                    # defaults to 'localhost'; use '*' for all
                    # (change requires restart)
#port = 5432                # (change requires restart)
max_connections = 100           # (change requires restart)
#superuser_reserved_connections = 3 # (change requires restart)
#unix_socket_directories = '/var/run/postgresql'    # comma-separated list of o
                    # (change requires restart)
...
```

This should be changed to

```
...
# - Connection Settings -

listen_addresses = ''      # what IP address(es) to listen on;
                    # comma-separated list Sof addresses;
                    # defaults to 'localhost'; use '*' for all
                    # (change requires restart)
#port = 5432                # (change requires restart)
max_connections = 100           # (change requires restart)
#superuser_reserved_connections = 3 # (change requires restart)
```

```
unix_socket_directories = 'DirectoryNameMatchingPGHOST'    # comma-separated li
                          # (change requires restart)
...
```

Very important notice: when we changed those two lines, we also removed the # from the start of the line. If you don't remove this, the line remains commented out and your changes will have no effect.

The string on the `unix_socket_directories` line must contain the same directory name as you get via the shell command:

```
$ echo $PGHOST
```

You can also change the value of `max_connections` to something smaller than 100 (e.g. 10). This is not necessary, but I do this to make PostgreSQL use slightly less runtime memory.

After making these changes, your PostgreSQL server is ready to go.

> One place where PostgreSQL is less space efficient than it might be is in the size of its transaction logs. These logs live in the directory `pgsql/data/pg_wal` and are essential for the functioning of your PostgreSQL server. If you remove any files from this directory, you will render your server inoperable. Similarly, manually changing the files under `pgsql/data/base` and its subdirectories will probably break your PostgreSQL server.
>
> If you mess up your PostgreSQL server badly enough, it will need to be re-installed. If such a thing happens, all of your databases are useless and all of the data in them is irretrievable. You will need to completely remove the `/localstorage/$USER/pgsql` directory and re-install as above. You should only do this in extreme circumstances; most people will install the server directories/files once, and they will use the same installation until the end of the term.
>
> If you need to remove the `pgsql` directory, then all of your databases and any data in them are gone forever. This is not a problem if you set up your databases by loading new views, functions, data, etc. from a file, but if you type `create` commands directly into the database, then the created objects will be lost. The best way to avoid such catastrophic loss of data is to type your SQL `create` statements into a file and load them into the database from there. Alternatively, you'd need to do regular back-ups of your databases using the `pg_dump` command.

Stage 4: Using your PostgreSQL Server

In this section, we assume that you have completed Stage 3 and now have a directory for PostgreSQL on `d.cse.unsw.edu.au`

When you want to do some work with PostgreSQL: login to `d.cse` set up your environment, start your server, do your work, and then stop the server before logging off

Do not leave your PostgreSQL server running while you are not using it.

A typical session with PostgreSQLwould begin with you logging in to the `d.cse.unsw.edu.au`. You would then do something like ...

```
$ source /localstorage/$USER/env
   ... sets up environment ...
$ p1
   ... start the PostgreSQL server ...
$ psql SomeDatabase
   ... work with a database ...
$ p0
   ... stop the PostgreSQL server ...
```

Each time you want to use your PostgreSQL server, you'll need to do the following:

Note that `p1` and `p0` are abbreviations defined in the `env`. They invoke the `pg_ctl` command which controls the operation of the PostgreSQL server.

After using `p1`, you can check whether your server is running via the command:

```
$ psql -l
```

Note: `l` is lower-case L, not the digit 1.

Try starting, checking, and stopping the server a few times.

Things occasionally go wrong, and knowing how to deal with them will save you lots of time. There's a discussion of common problems at the end of this document; make sure that you read and understand it.

Once your PostgreSQL server is running, you can access your PostgreSQL databases via the `psql` command. You normally invoke this command by specifying the name of a database, e.g.

```
$ psql MyDatabase
```

If you type `psql` command without any arguments, it assumes that you are trying to access a database with the same name as your login name. Since you probably won't have created such a database, you're likely to get a message like:

```
psql: FATAL:  database "${USER}" does not exist
```

You will get a message like this any time that you try to access a database that does not exist.

If you're not sure what databases you have created, `psql` can tell you via the `-l` option

```
$ psql -l
```

If you run this command now, you ought to see output that looks like:

```
SET
                         List of databases
   Name    | Owner | Encoding |  Collate    |    Ctype    | Access privileges
-----------+-------+----------+-------------+-------------+-------------------
 postgres  | jas   | UTF8     | C           | en_AU.UTF-8 |
 template0 | jas   | UTF8     | C           | en_AU.UTF-8 | =c/jas           +
           |       |          |             |             | jas=CTc/jas
 template1 | jas   | UTF8     | en_US.utf8  | en_US.utf8  |
(3 rows)
```

Of course, it will be *your* username, and not `jas`.

Note that PostgreSQL commands like `psql` and `createdb` are a lot noisier than normal Linux commands. In particular, they all seem to print `SET` when they run; you can ignore this. Similarly, if you see output like `INSERT 0 1`, you can ignore that as well.

The above three databases are created for use by the PostgreSQL server; you should not modify them. At this stage, you don't need to worry about the contents of the other columns in the output. As long as you see at least three databases when you run the `psql -l` command, it means that your PostgreSQL server is up and running ok.

Note that you are the administrator for your PostgreSQL server (add "database administrator" to your CV) and you can create as many databases as you like, within the limits of your disk quota.

From within `psql`, the fact that you are an administrator is indicated by a prompt that looks like

```
dbName=#
```

rather than the prompt for database users

```
dbName=>
```

which you may have seen in textbooks or notes.

Note that you can only access databases created as above while you're logged into `d.cse`. In other words, you must run the `psql` command on `d.cse`.

Note that the **only** commands that you should run on `db.cse` are the commands to start and stop the server, the `psql` command to start an interactive session with a database, and the other PostgreSQL clients such as `createdb`. Do not run other processes such as web browsers or drawing programs or games on `d.cse`. Text editors are OK; VScode is not.

All of the PostgreSQL client applications are documented in the PostgreSQL manual, in the "PostgreSQL Client Applications" section. While there are quite a few PostgreSQL client commands, `psql` will be the one that you will mostly use.

> **Mini-Exercise:** a quick way to check whether your PostgreSQL server is running is to try the command:
>
> ```
> $ psql -l
> ```
>
> Try this command now.
>
> If you get a response like:
>
> ```
> psql: command not found
> ```
>
> then you haven't set up your environment properly; `source` the `env` file.
>
> If you get a response like:
>
> ```
> psql: could not connect to server: No such file or directory
>         Is the server running locally and accepting
>         connections on Unix domain socket "....s.PGSQL.5432"?
> ```
>
> then the server isn't running.

> If you get a list of databases, like the example above, then this means your server is running ok and ready for use.

Cleaning up

After you've finished a session with PostgreSQL, it's essential that you shut your PostgreSQL server down (to prevent overloading `d.cse`). Recall that your this via the command

```
$ p0
```

PostgreSQL generates log files that can potentially grow quite large. If you start your server using `p1`, the log file is called

```
/localstorage/$USER/pgsql/data/log
```

It would be worth checking every so often to see how large it has become. To clean up the log, simply stop the server and remove the file. Note: if you remove the logfile while the server is running, you may not remove it at all; its link in the filesystem will be gone, but the disk space will continue to be used and grow until the server stops.

> **Mini-Exercise:** Try starting and stopping the server a few times, and running `psql` both when the server is running and when it's not, just to see the kinds of messages you'll get.

Exercise #1: Making a database

Start by logging in to `d.cse` and setting the environment.

Once the PostgreSQL server is running, try creating a database by running the command:

```
$ createdb mydb
```

which will create the database, or give an error message if it can't create it for some reason. (A typical reason for failure would be that your PostgreSQL server is not running.)

Now use the `psql -l` command to check that the new database exists.

You can access the database by running the command:

```
$ psql mydb
```

which should give you a message like

```
SET
psql (13.3 (Debian 13.3-1))
Type "help" for help.

mydb=#
```

Note that `psql` lets you execute two kinds of commands: SQL queries and updates, and `psql` "meta"-commands. The `psql` "meta"-commands allow you to examine the database schema, and control various aspects of `psql` itself, such as where it writes its output and how it formats tables.

Getting back to the `psql` session that you just started, the `mydb` database is empty, so there's not much you can do with it. The `\d` (describe) command allows you to check what's in the database. If you type it now, you get the unsurprising response

```
mydb=# \d
No relations found.
```

About the only useful thing you can do at the moment is to quit from `psql` via the `\q` command.

```
mydb=# \q
$ ... now waiting for you to type Linux commands ...
```

Note: it is common to forget which prompt you're looking at and sometimes type Linux commands to `psql` or to type SQL queries to the Linux shell. It usually becomes apparent fairly quickly what you've done wrong, but can initially be confusing when you think that the command/query is not behaving as it should. Here are examples of making the above two mistakes:

```
$ ... Linux command interpreter ...
$ select * from table;
-bash: syntax error near unexpected token `from'
$ psql mydb
... change context to PostgreSQL ...
mydb=# ls -l
mydb-# ... PostgreSQL waits for you to complete what it thinks is an SQL query
mydb-# ;    ... because semi-colon finishes and then executes an SQL query ...
ERROR:  syntax error at or near "ls"
LINE 1: ls -l
        ^
mydb=# \q
$ ... back to Linux command interpreter ...
```

Exercise #2: Populating a database

Once the `mydb` database exists, the following command will create the schemas (tables) and populate them with tuples:

```
$ psql mydb -f /home/cs3311/web/21T3/pracs/02/mydb.sql
```

Note that this command produces quite a bit of output, telling you what changes it's making to the database. The output should look like:

```
SET
CREATE TABLE
INSERT 0 1
INSERT 0 1
INSERT 0 1
CREATE TABLE
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
CREATE TABLE
```

```
INSERT 0 1
INSERT 0 1
INSERT 0 1
CREATE TABLE
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
INSERT 0 1
```

The lines containing `CREATE TABLE` are, obviously, related to PostgreSQL creating new database tables (there are four of them). The lines containing `INSERT` are related to PostgreSQL adding new tuples into those tables.

Clearly, if we were adding hundreds of tuples to the tables, the output would be very long. You can get PostgreSQL to stop giving you the `INSERT` messages by using the `-q` option to the `psql` command.

PostgreSQL's output can be verbose during database loading. If you want to ignore everything *except* error messages, you could use a command like:

```
$ psql mydb -f /home/cs3311/web/21T3/pracs/02/mydb.sql 2>&1 | grep ERROR
```

If you don't understand the fine details of the above, take a look at the documentation for the Linux/Unix shell.

The `-f` option to `psql` tells it to read its input from a file, rather than from standard input (normally, the keyboard). If you look in the `mydb.sql` file, you'll find a mix of table (relation) definitions and statements to insert tuples into the database. We don't expect you to understand the contents of the file at this stage.

If you try to run the above command again, you will generate a heap of error messages, because you're trying to insert the same collection of tables and tuples into the database, when they've already been inserted.

Note that the tables and tuples are now permanently stored on disk. If you switch your PostgreSQL server off, when you restart it the contents of the `mydb` database will be available, in whatever state you left them from the last time you used the database.

Exercise #3: Examining a database

One simple way to manipulate PostgreSQL databases is to use the `psql` command (which is a "shell" like the `sqlite3` command in the first prac exercise). A useful way to start exploring a database is to find out what tables it has. We saw before that you can do this with the \d (describe) command. Let's try that on the newly-populated `mydb` database.

```
mydb=# \d
         List of relations
 Schema |    Name    | Type  | Owner
--------+------------+-------+-------
```

```
 public | courses   | table | $USER
 public | enrolment | table | $USER
 public | staff     | table | $USER
 public | students  | table | $USER
(4 rows)
```

You can ignore the `Schema` column for the time being. The `Name` column tells you the names of all tables (relations) in the current database instance. The `Type` column is obvious, and, you may think, unnecessary. It's there because `\d` will list all objects in the database, not just tables; it just happens that there are only tables in this simple database. The `Owner` should be your username, for all tables.

One thing to notice is that the table names are all in lower-case, whereas in the `mydb.sql` file, they had an initial upper-case letter. The SQL standard says that case does not matter in unquoted identifiers and so `Staff` and `staff` and `STAFF` and even `StAfF` are all equivalent. To deal with this, PostgreSQL simply maps *identifiers* into all lower case internally. You can still use `Staff` when you're typing in SQL commands; it will be mapped automatically before use.

> There are, however, advantages to using all lower case whenever you're dealing with `psql`. For one thing, it means that you don't have to keep looking for the shift-key. More importantly, `psql` provides table name and field name completion (you type an initial part of a table name, then type the TAB key, and `psql` completes the name for you if it has sufficient context to determine this unambiguously), but it only works when you type everything in lower case. The `psql` interface has a number of other features (e.g. history, command line editing) that make it very nice to use.

If you want to find out more details about an individual table, you can use:

```
mydb=# \d Staff
           Table "public.staff"
  Column   |         Type          | Modifiers
-----------+-----------------------+-----------
 userid    | character varying(10) | not null
 name      | character varying(30) |
 position  | character varying(20) |
 phone     | integer               |
Indexes:
    "staff_pkey" PRIMARY KEY, btree (userid)
Referenced by:
    TABLE "courses" CONSTRAINT "courses_lecturer_fkey" FOREIGN KEY (lecturer) F
```

As you can see, the complete name of the table is `public.staff`, which includes the schema name. PostgreSQL has the notion of a "current schema" (which is the schema called `public`, by default), and you can abbreviate table names by omitting the current schema name, which is what we normally do. The types of each column look slightly different to what's in the `mydb.sql` file; these are just PostgreSQL's internal names for the standard SQL types in the schema file. You can also see that the `userid` field is not allowed to be null; this is because it's the primary key (as you can see from the index description) and primary keys may not contain null values. The index description also tells you that PostgreSQL has built a B-tree index on the `userid` field.

The final line in the output tells you that one of the other tables in the database (`Courses`) has a foreign key that refers to the primary key of the `Staff` table, which you can easily see by looking at

the `mydb.sql` file. This is slightly useful for a small database, but becomes extremely useful for larger databases with many tables.

The next thing we want to find out is what data is actually contained in the tables. This requires us to use the SQL query language, which you may not know yet, so we'll briefly explain the SQL statements that we're using, as we do them.

We could find out all the details of staff members as follows:

```
mydb=# select * from Staff;
  userid  |     name      |    position     | phone
----------+---------------+-----------------+-------
 jingling | Jingling Xue  | Professor       | 54889
 jas      | John Shepherd | Senior Lecturer | 56494
 andrewt  | Andrew Taylor | Senior Lecturer | 55525
(3 rows)
```

The SQL statement says, more or less, "tell me everything (*) about the contents of the `Staff` table". Each row in the output below the heading represents a tuple in the table.

Note that the SQL statement ends with a semi-colon. The meta-commands that we've seen previously didn't require this, but SQL statements can be quite large, and so, to allow you to type them over several lines, the system requires you to type a semi-colon to mark the end of the SQL statement.

If you forget to put a semi-colon, the prompt changes subtly:

```
mydb=# select * from Staff
mydb-#
```

This is PostgreSQL's way of telling you that you're in the middle of an SQL statement and that you'll eventually need to type a semi-colon. If you then simply type a semi-colon to the second prompt, the SQL statement will execute as above.

Mini-Exercise: find out the contents of the other tables.
Here are some other SQL statements for you to try out. You don't need to understand their structure yet, but they'll give you an idea of the kind of capabilities that the SQL language offers.

- Which students are studying for a CS degree (3778)?

  ```
  select * from Students where degree=3778;
  ```

- How many students are studying for a CS degree?

  ```
  select count(*) from Students where degree=3778;
  ```

- Who are the professors?

  ```
  select * from Staff where position ilike '%professor%';
  ```

- How many students are enrolled in each course?

  ```
  select course,count(*) from Enrolment group by course;
  ```

- Which courses is Andrew Taylor teaching?

  ```
  select c.code, c.title
  from   Courses c, Staff s
  ```

```
            where  s.name='Andrew Taylor' and c.lecturer=s.userid;
```

or

```
select c.code, c.title
from   Courses c join Staff s on (c.lecturer=s.userid)
where  s.name='Andrew Taylor';
```

The last query is laid out as we normally lay out more complex SQL statements: with a keyword starting each line, and each clause of the SQL statement starting on a separate line.

Try experimenting with variations of the above queries.

Sorting out Problems

It is very difficult to diagnose problems with software over email, unless you give sufficient details about the problem. An email that's as vague as "My PostgreSQL server isn't working. What should I do?", is basically useless. Any email about problems with software should contain details of

- what you were attempting to do
- precisely what commands you used
- exactly what output you got

One way to achieve this is to copy-and-paste the last few commands and responses into your email.

Alternatively, you should come to a consultation where we can work through the problem via screen sharing (which is usually very quick).

Can't shut server down?

When you use `p0` to shut down your PostgreSQL server, you'll observe something like:

```
$ p0
waiting for server to shut down....
```

Dots will keep coming until the server is finally shut down, at which point you will see:

```
$ p0
waiting for server to shut down........ done
server stopped
```

Sometimes, you'll end up waiting for a long time and the server still doesn't shut down. This is typically because you have an `psql` session running in some other window (the PostgreSQL server won't shut down until all clients have disconnected from the server). The way to fix this is to find the `psql` session and end it. If you can find the window where it's running, simply use `\q` to quit from `psql`. If you can't find the window, or it's running from a different machine (e.g. you're in the lab and find that you left a `psql` running at home), then use `ps` to find the process id of the `psql` session and stop it using the Linux `kill` command.

Can't restart server?

Occasionally, you'll find that your PostgreSQL server was not shut down cleanly the last time you used it and you cannot re-start it next time you try to use it. We'll discuss how to solve that here ...

The typical symptoms of this problem are that you log in to `d.cse`, set up your environment, try to start your PostgreSQL server and you get the message:

```
pg_ctl: another server may be running; trying to start server anyway
waiting for server to start.... stopped waiting
pg_ctl: could not start server
Examine the log output.
```

When you go and check the log file, you'll probably find, right at the end, something like:

```
$ tail -2 $PGDATA/Log
FATAL:  lock file "postmaster.pid" already exists
HINT:  Is another postmaster (PID NNNN) running in data directory "/localstorag
```

where *NNNN* is a number.

There are two possible causes for this: the server is already running, or the server did not terminate properly after the last time you used it. You can check whether the server is currently running by the command `psql -l`. If that gives you a list of your databases, then you simply forgot to shut the server down last time you used it and it's ready for you to use again. If `psql -l` tells you that there's no server running, then you'll need to do some cleaning up before you can restart the server ...

When the PostgreSQL server is run, it keeps a record of the Unix process that it's running as in a file called:

```
$PGDATA/postmaster.pid
```

Normally when your PostgreSQL server process terminates (e.g. via `p0`), this file will be removed. If your PostgreSQL server stops, and this file persists, then `p1` becomes confused and thinks that there is still a PostgreSQL server running even though there isn't.

The first step in cleaning up is to remove this file:

```
$ rm $PGDATA/postmaster.pid
```

You should also clean up the socket files used by the PostgreSQL server. You can do this via the command:

```
$ rm $PGDATA/.s*
```

Once you've cleaned all of this up, then the `p1` command ought to allow you to start your PostgreSQL server ok.

Happy PostgreSQL'ing, *jas*